



AI Infrastructure: Storage Options



Congratulations on completing the course. This course summary is your review guide. Print it for a handy reference.




The data preparation phase addresses two key business metrics:

- Model accuracy
- Model freshness.


The following storage options help during the prepare and train phase:

- Cloud Storage Fuse
- Managed Lustre


Cloud Storage FUSE is a tool that helps you use Cloud Storage as if it were a local file system.




Google Cloud provides Anywhere Cache, an SSD-based zonal read cache built directly on top of Cloud Storage.



Managed Lustre is a distributed parallel file system, which means it's built to enable thousands of compute nodes to read and write to a single file system simultaneously.




AI Infrastructure: Storage Options




Congratulations on completing the course. This course summary is your review guide. Print it for a handy reference.


The ideal storage solution for inference depends on your update frequency, node count, and performance needs.



When using Cloud Storage as the source of truth, there are two main storage options to retrieve your data at the GKE-pod level: Cloud Storage Fuse and Hyperdisk ML.




GKE Image Streaming allows GKE to start containers without waiting for the entire image to be downloaded. Instead of pulling all layers, GKE streams the image data from the registry on demand as your application needs it.



Hyperdisk ML disks are only available in a single zone. Optionally, you can use the Hyperdisk ML multi-zone feature to dynamically link multiple zonal disks (that contain the same content) in a single logical PersistentVolumeClaim and PersistentVolume.

To optimize costs, Cloud Storage offers different storage classes based on how often you anticipate accessing your archived data. The less often you plan to access the data, the lower the storage cost.



Additional Resources

- [Caching in Cloud Storage FUSE](#)
- [AI Hypercomputer Storage services](#)
- [Storage for AI and ML workloads](#)